

Classification

Using Machine Learning Tools

Geron Chapter 3, 6

Last time ...

- Steps you can take to improve the performance of your ML system:
 - replacing data
 - scaling data
- Pipelines for repeatable workflows
- Linear regression model
- Polynomial models
- Parameters and hyperparameters

Classification terminology

- **Classifier**

- Uses a **discrete set** of possible outputs = classes
- Can be supervised, semi-supervised or unsupervised (see week 7)

- **Target**

- What we are predicting
- Passed in as “y” to regression/classification method
- Also called: **label, ground truth, dependent variable, outcome variable, or response variable**

Regression vs Classification

- Regression:
 - predicts real numbers (values)
 - on a numerical, ordered scale
 - the larger the difference, the worse
 - e.g. house price, wine quality
- Classification:
 - predicts classes (labels)
 - typically categorical, normally no meaningful order
 - all differences are usually treated equally
 - e.g. cancer vs healthy ; kangaroo vs pademelon vs quokka

Linear Model as a Binary Classifier

- Linear model used in regression:

$$\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

- Apply a threshold:

If $\hat{y} < \text{threshold}$

Class 1

Else

Class 2

Stochastic Gradient Descent Classifier

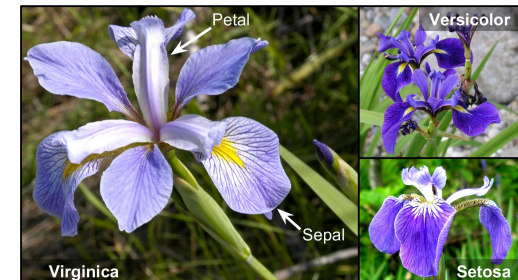
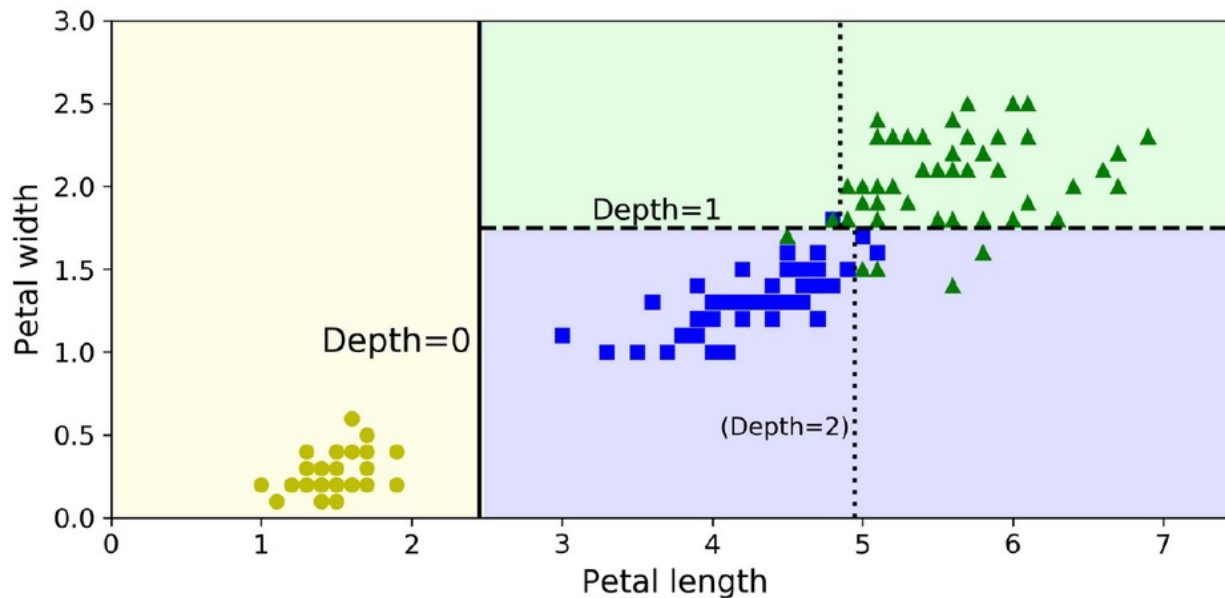
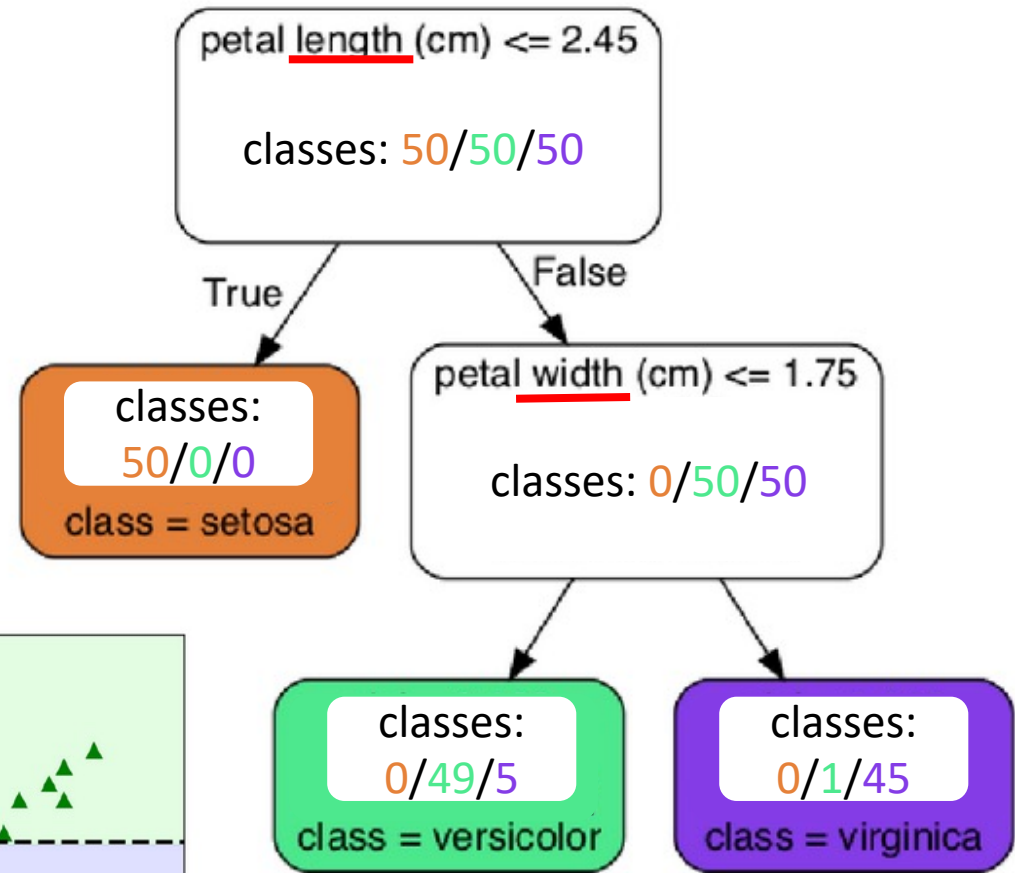
- Binary classifier using a linear model
- Stochastic Gradient Descent (SGD) is a fitting algorithm
 - Iteratively follows the gradient (derivative) of the loss function
 - Fast, scalable
- SGD classifier in scikit-learn is a linear model using SGD

```
from sklearn.linear_model import SGDClassifier  
clf = SGDClassifier(random_state=42)  
clf.fit(X_train, y_train)
```

Decision Tree

Images: Geron, Hands On ML

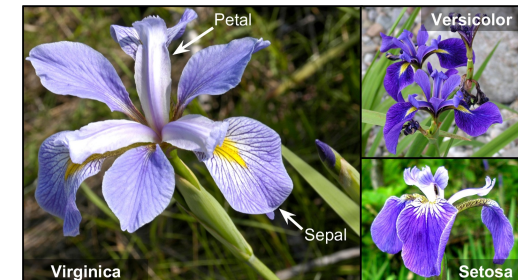
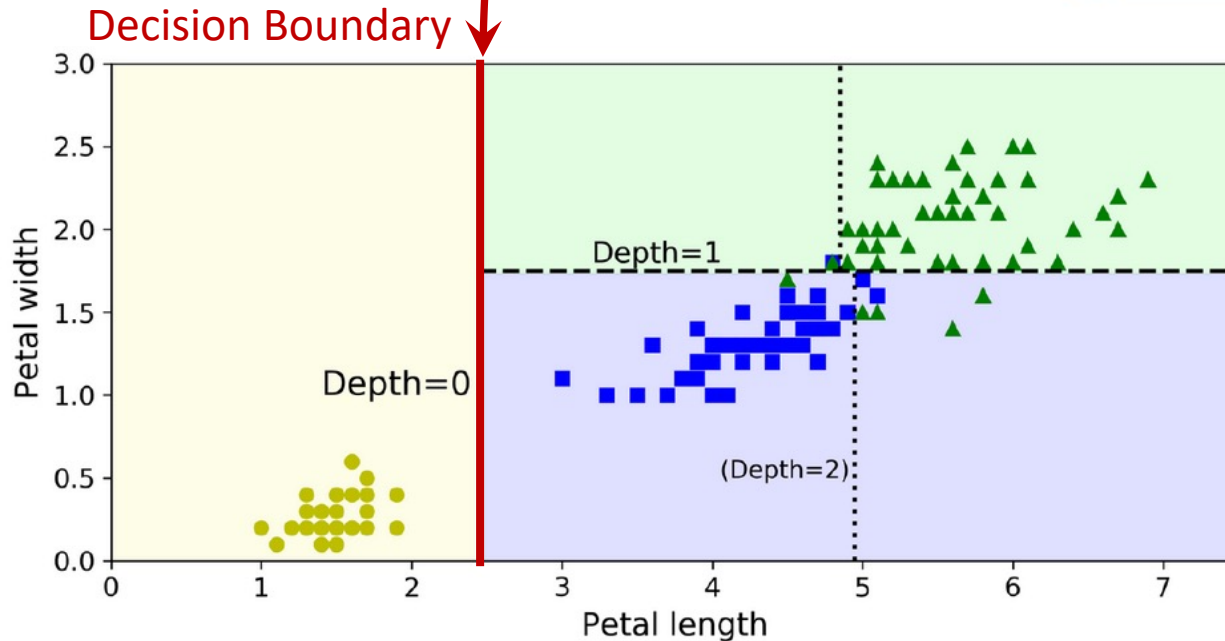
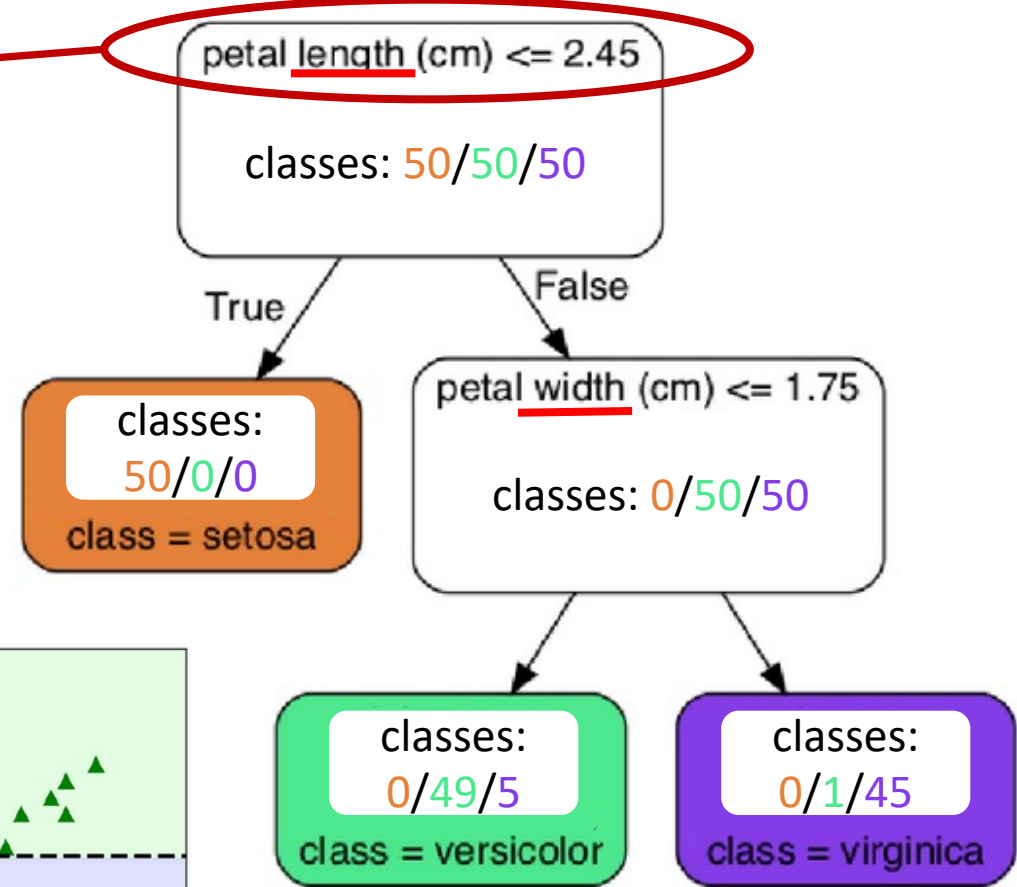
- Iterative splitting
- Maximise class separation
- One feature at a time
- Up to maximum depth
- Prune to avoid overfitting



Decision Tree

Images: Geron, Hands On ML

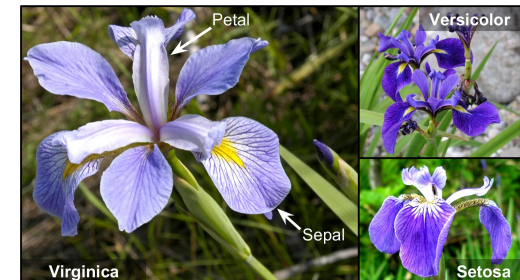
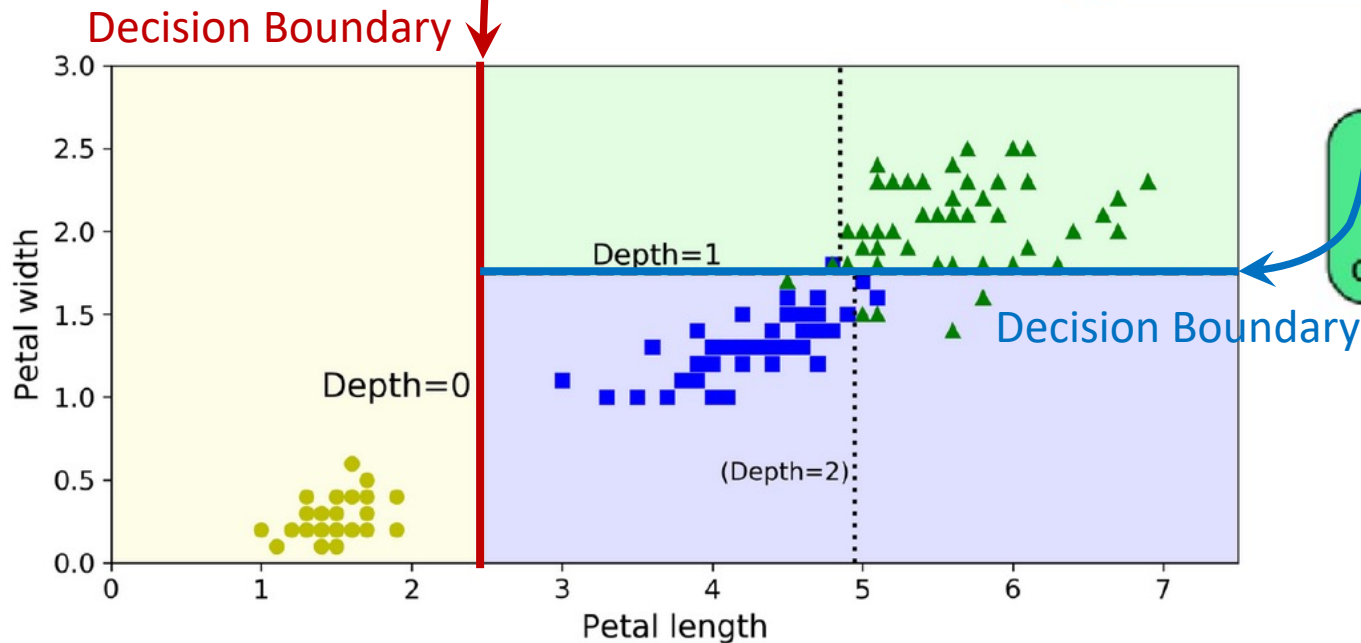
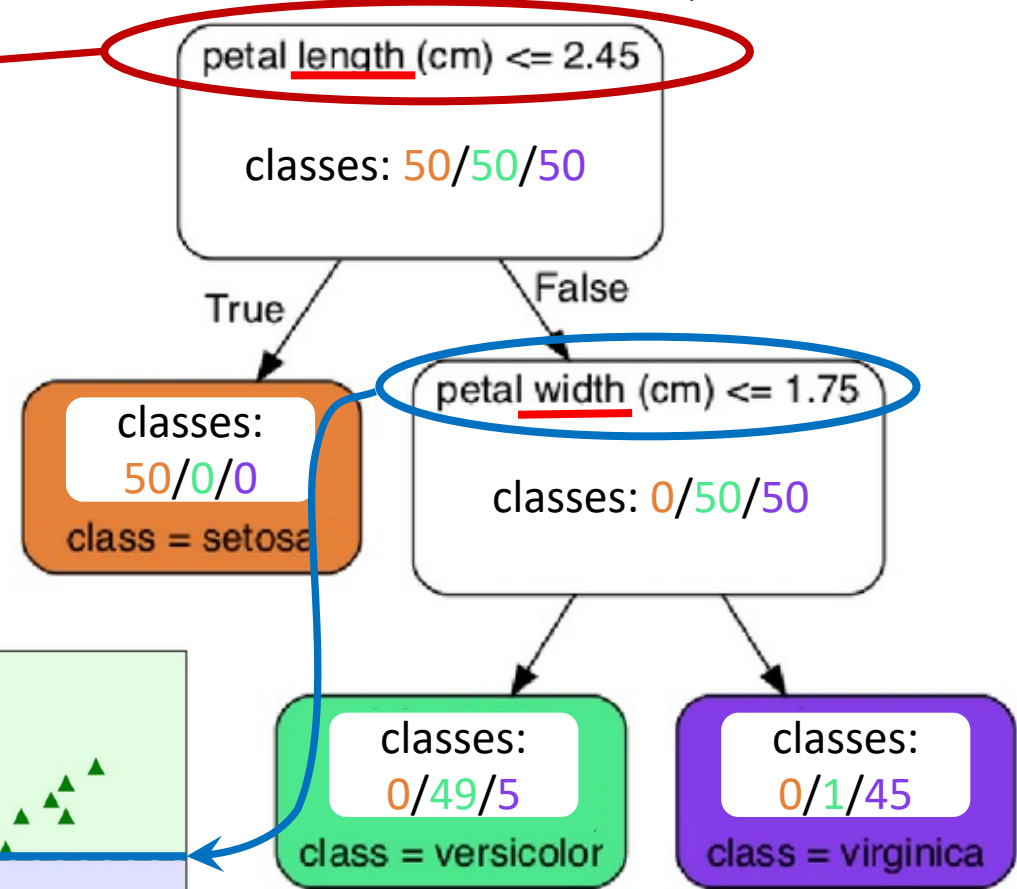
- Iterative splitting
- Maximise class separation
- One feature at a time
- Up to maximum depth
- Prune to avoid overfitting



Decision Tree

Images: Geron, Hands On ML

- Iterative splitting
- Maximise class separation
- One feature at a time
- Up to maximum depth
- Prune to avoid overfitting



Inner workings → white-box

$$G_i = 1 - \sum_{k=1}^n p_{i,k}^2$$

- Gini impurity metric measures the class distribution in a node
 - best = only one class (pure) → $G = 0$ for impurity
 - worst = completely evenly spread → $G_{\max} = 1 - 1/N$ for impurity

- CART = Classification and Regression Tree

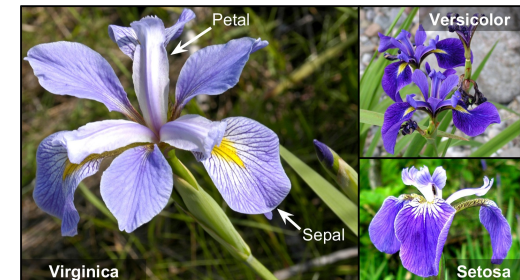
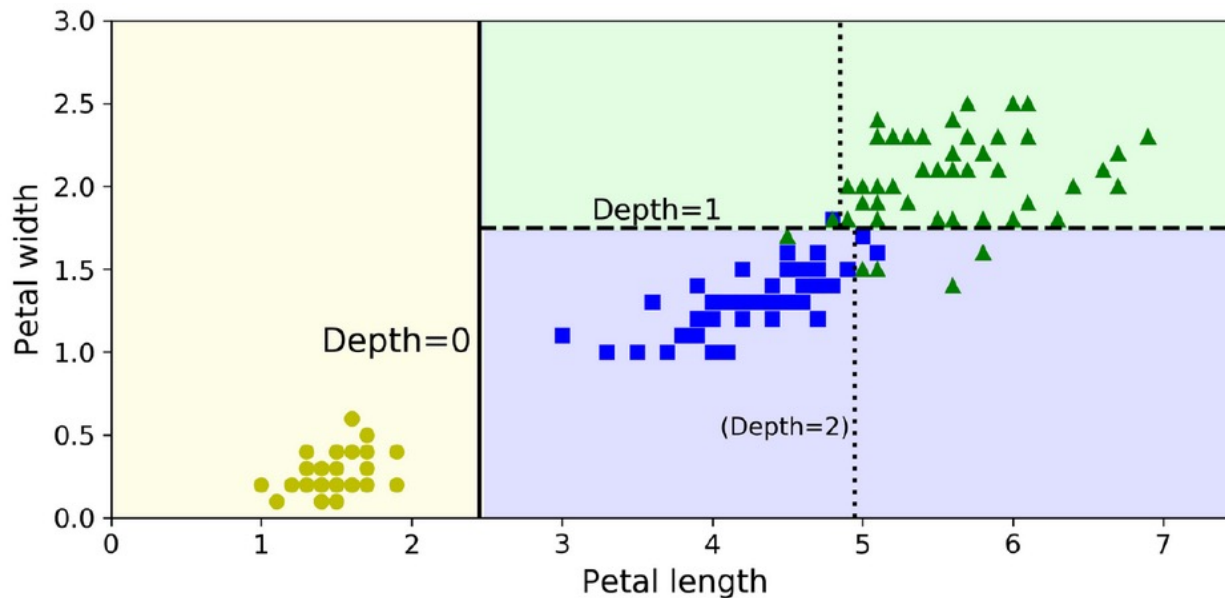
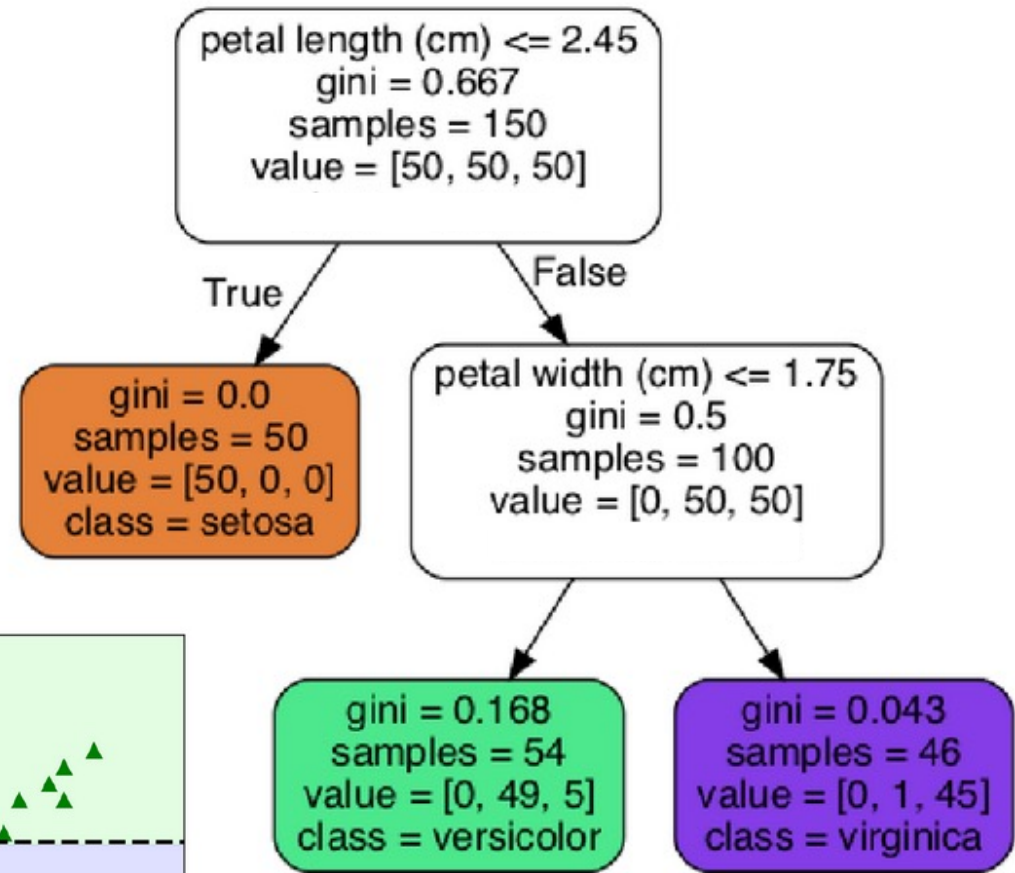
$$J(k, t_k) = \frac{m_{\text{left}}}{m} G_{\text{left}} + \frac{m_{\text{right}}}{m} G_{\text{right}}$$

- A loss function is minimised to set threshold → min impurity
 - it is based on a weighted sum of the impurities from each branch
- Search for best threshold and branch out
- Iteratively grow tree until maximum depth reached

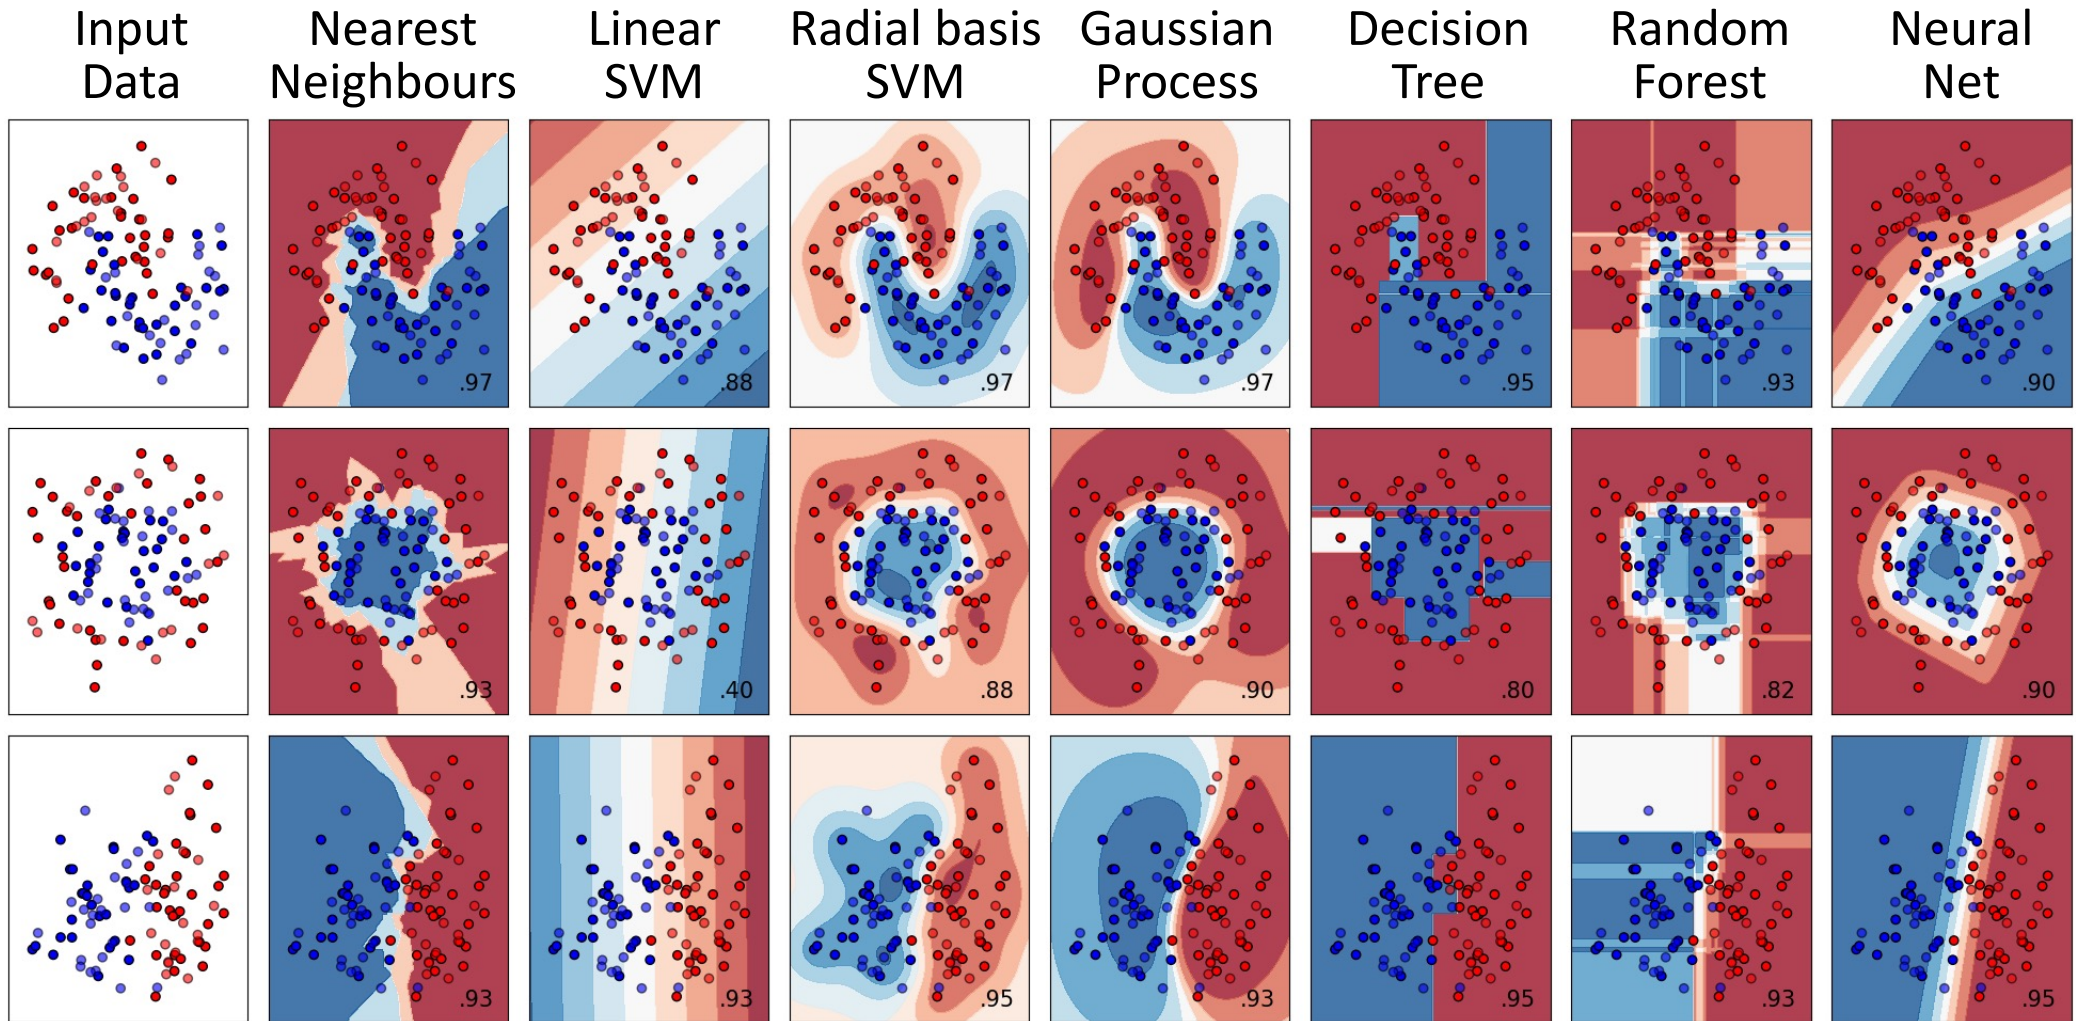
Decision Tree

Images: Geron, Hands On ML

- Iterative splitting
- Maximise class separation
- One feature at a time
- Up to maximum depth
- Prune to avoid overfitting



Comparison of Classification Approaches



Note the different
decision boundary shapes

Image: © 2007 - 2019, scikit-learn developers
(BSD 3-clause License)

Evaluation / Performance Metrics

Evaluation / Performance Metrics

- Regression:

- (Root) Mean Squared Error = L^2 norm = $\frac{1}{N} \sum_i |x_i - y_i|^2 = \|\mathbf{x} - \mathbf{y}\|_2$

- Mean Absolute Error = L^1 norm = $\frac{1}{N} \sum_i |x_i - y_i| = \|\mathbf{x} - \mathbf{y}\|_1$

- Median/Max Absolute Error

- R^2 / correlation coefficient $\rho = r = \frac{(\mathbf{x} - \bar{x}) \cdot (\mathbf{y} - \bar{y})}{|\mathbf{x} - \bar{x}| |\mathbf{y} - \bar{y}|} = \cos \theta$

Evaluation / Performance Metrics

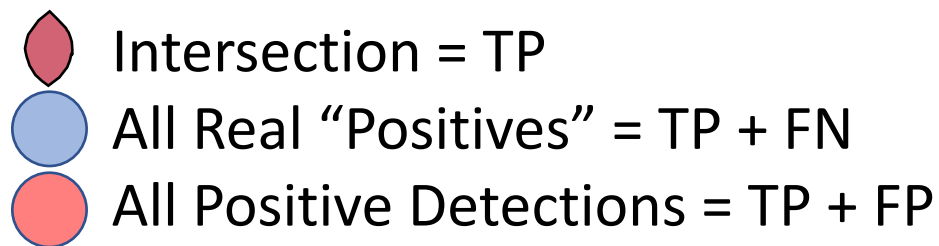
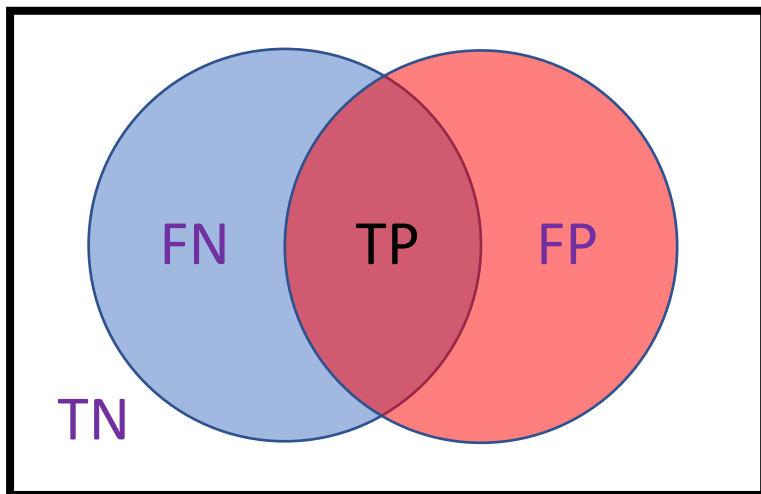
- Regression:

- (Root) Mean Squared Error = L^2 norm = $\frac{1}{N} \sum_i |x_i - y_i|^2 = \|\mathbf{x} - \mathbf{y}\|_2$
- Mean Absolute Error = L^1 norm = $\frac{1}{N} \sum_i |x_i - y_i| = \|\mathbf{x} - \mathbf{y}\|_1$
- Median/Max Absolute Error
- R^2 / correlation coefficient $\rho = r = \frac{(\mathbf{x} - \bar{x}) \cdot (\mathbf{y} - \bar{y})}{|\mathbf{x} - \bar{x}| |\mathbf{y} - \bar{y}|} = \cos \theta$

- Classification:

- Accuracy / Confusion Matrix
- Precision / Recall
- ROC Curve / Area Under Curve (AUC)
- F_1 Score

Positive & Negative / True & False



- **Positives/Negatives** (P/N) are from **model predictions** (above/below threshold)
- **True** is when the prediction **agrees with the ground truth** (labels)
- A true positive is a positive that agrees with ground truth
- A **false negative** is something of interest in the ground truth that is **missed** by the model
- Both FP and FN are bad, but usually not equally bad
- Sometimes P/N is obvious (e.g. detecting an object) but sometimes it is arbitrary (e.g. male/female) or even counter-intuitive (e.g. disease = positive)

Precision and Recall

$$Precision = \frac{TP}{TP + FP} = \frac{\text{Intersection}}{\text{All Positive Detections}}$$

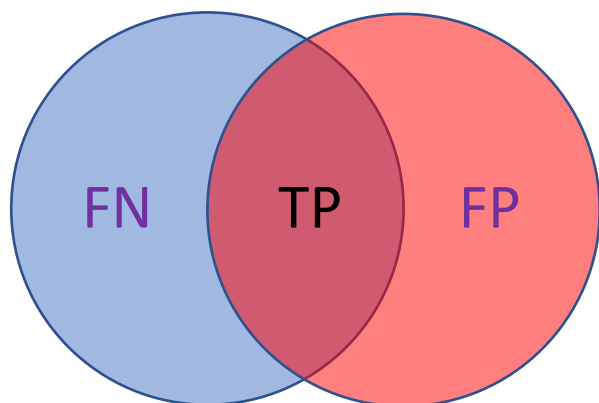
What fraction of positive predictions are correct?




= Positive predictive value

$$Recall = \frac{TP}{TP + FN} = \frac{\text{Intersection}}{\text{All Real "Positives"}}$$

What fraction of the real positive class are detected?

= True positive rate
= Sensitivity



-  Intersection = TP
-  All Real "Positives" = TP + FN
-  All Positive Detections = TP + FP

Precision and Recall

$$\textit{Precision} = \frac{TP}{TP + FP}$$

What fraction of positive predictions are correct?

= Positive predictive value

$$\textit{Recall} = \frac{TP}{TP + FN}$$

What fraction of the real positive class are detected?

= True positive rate

= Sensitivity

- Both ignore True Negatives (TN)
- Compare with Accuracy = $(TP+TN)/(TP+TN+FP+FN)$
- Sometimes accuracy is better, sometimes not

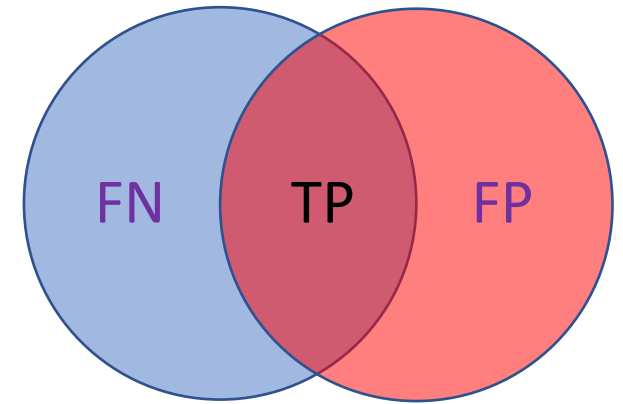
```
from sklearn.metrics import precision_score, recall_score
precision_score(y_train, y_train_pred)
recall_score(y_train, y_train_pred)
```

F₁ Score

$$F_1 = \frac{2}{\frac{1}{precision} + \frac{1}{recall}} = \frac{precision \times recall}{\frac{1}{2} * (precision + recall)} = \frac{TP}{\frac{1}{2} * (FN + FP + 2 * TP)}$$

$$= \frac{\text{Intersection}}{\frac{1}{2} \left[\text{All Real Positives} + \text{All Positive Detections} \right]}$$

- Harmonic mean of precision and recall
- Ignores True Negatives (TN)
- Evenly weights precision and recall
- F_β weights differently, using β value

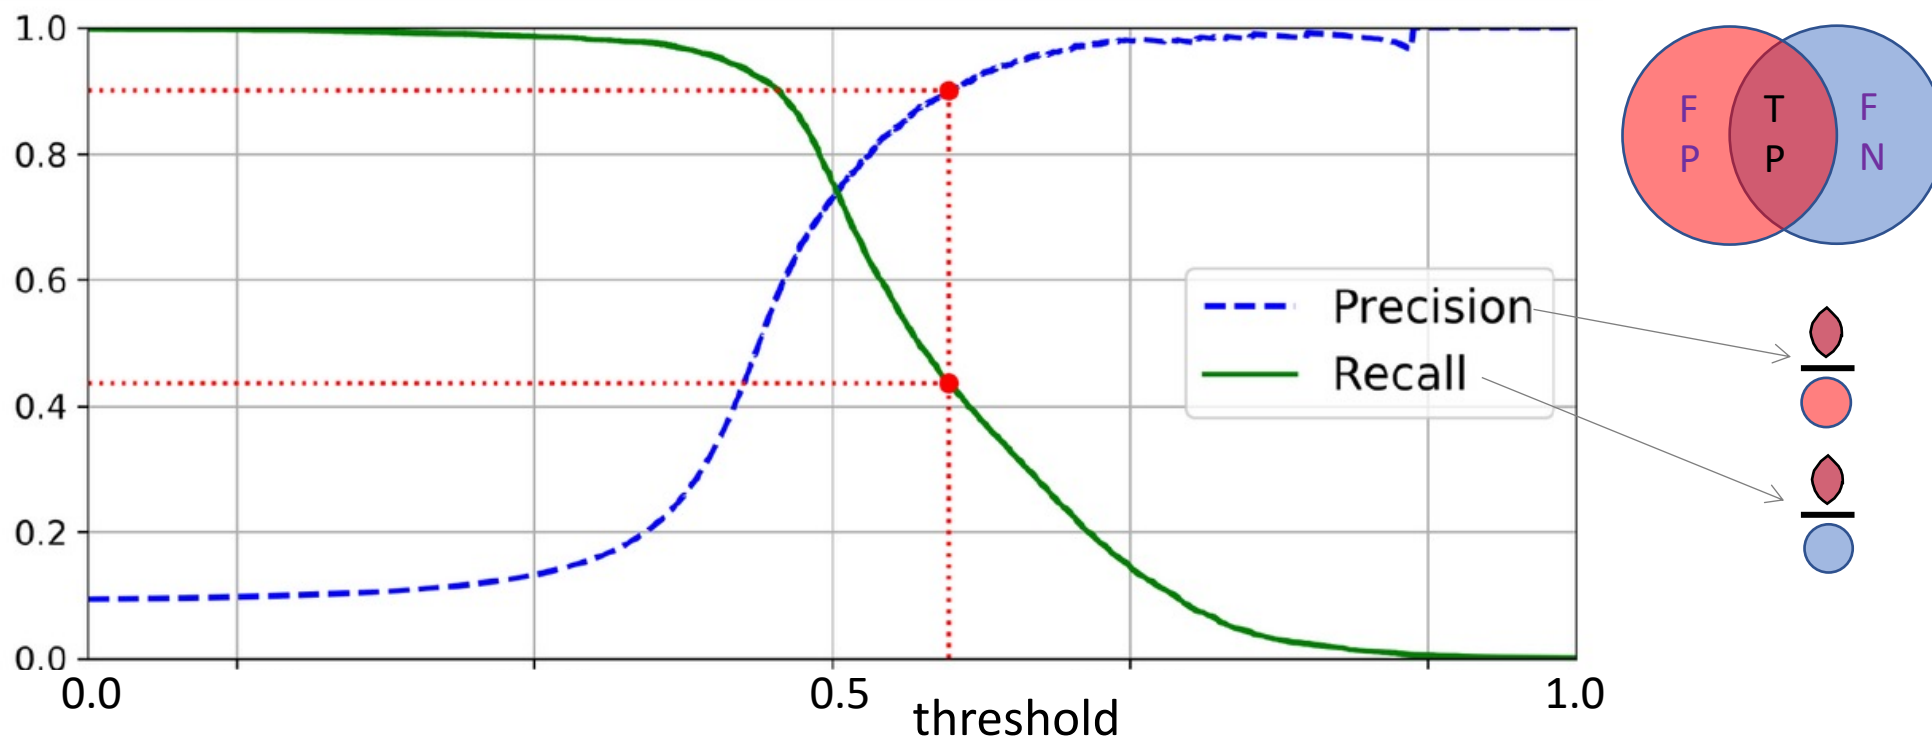


- Intersection = TP
- All Real Positives = TP + FN
- All Positive Detections = TP + FP

```
from sklearn.metrics import f1_score  
f1_score(y_train, y_train_pred)
```

Precision – Recall Trade-Off

Take real output value and **threshold** it \rightarrow TP, FP, FN, TN



Low threshold \rightarrow everything is positive, $FN=0$, $\text{blue circle} = \text{red diamond}$, Recall $\rightarrow 1$

High threshold \rightarrow everything is negative, $FP=0$, $\text{red circle} = \text{red diamond}$, Prec. $\rightarrow 1$

Image: Geron, Hands On ML

Receiver-Operating-Characteristic (ROC)

- Display trade-off between true positive rate and false positive rate
- Top left corner is best: TP=1, FP=0
- Each threshold \rightarrow one point
- Use all thresholds of the output to get curve (min \rightarrow max)
- Output of method (before thresholding) can be called a “probability”
- Can also be used to look at hyper-parameter changes (instead of threshold)

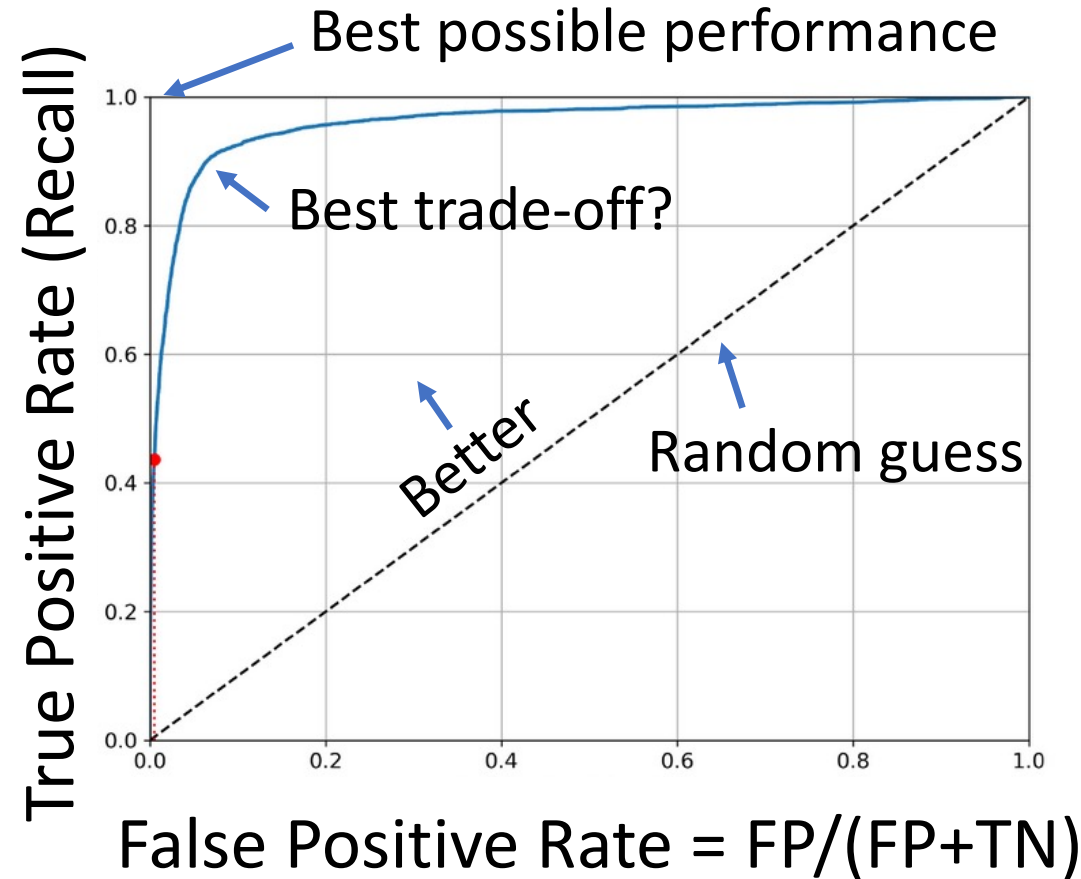


Image: Geron, Hands On ML,
added annotations

Area-Under-Curve (AUC)

- Single scalar value to compare classifiers or hyperparameters
- Depends on performance across all thresholds or hyperparameter settings
- Usually use whole curve

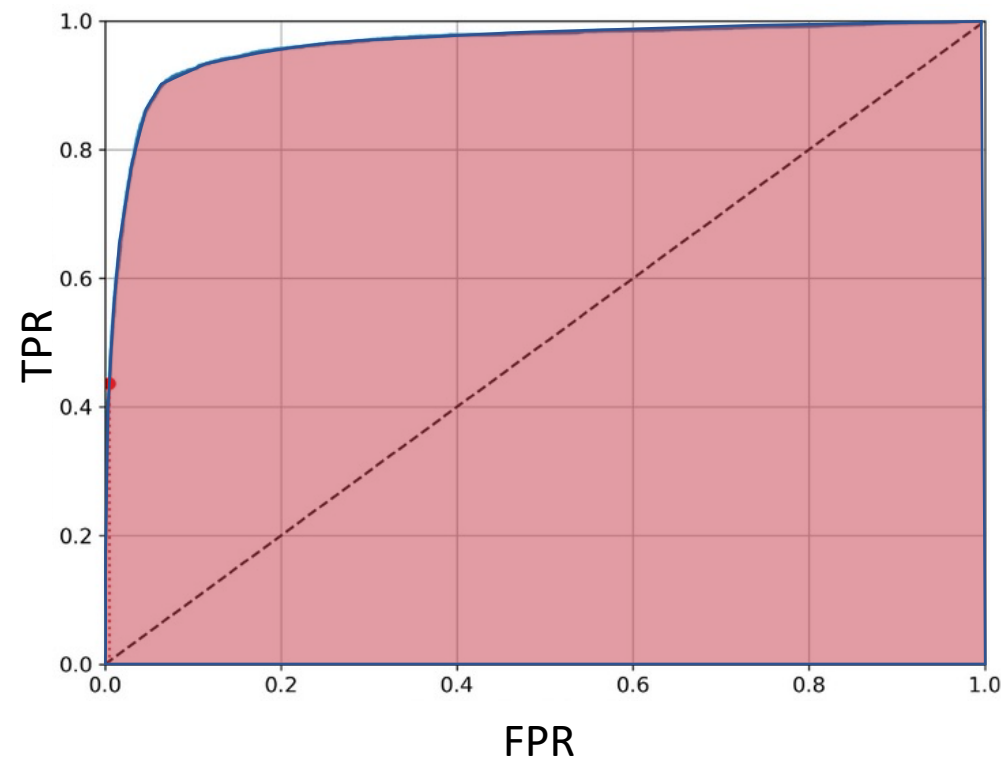


Image: Geron, Hands On ML,
added annotations

Area-Under-Curve (AUC)

- Single scalar value to compare classifiers or hyperparameters
- Depends on performance across all thresholds or hyperparameter settings
- Usually use whole curve
 - but a smaller range of FPR can also be used (based on acceptable performance range)

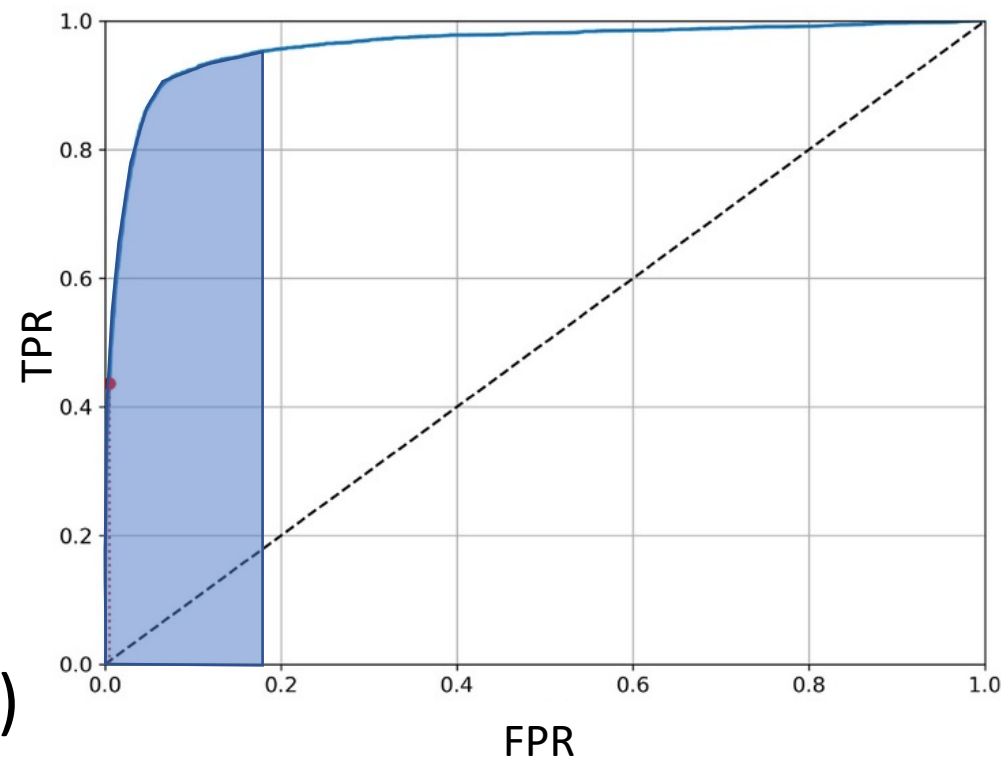
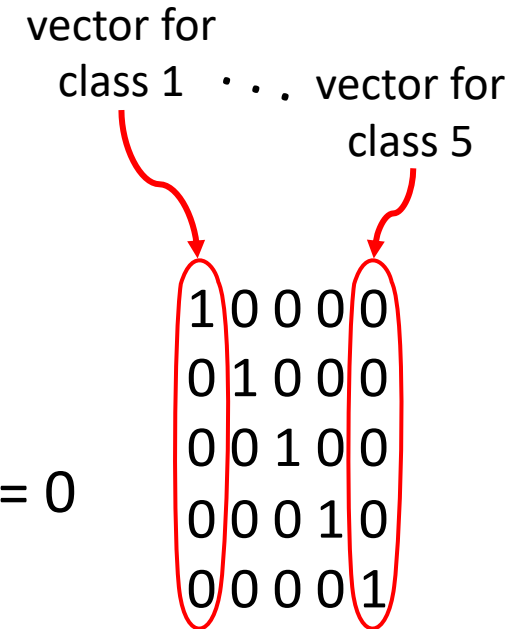


Image: Geron, Hands On ML,
added annotations

Multiclass Classification

- More than two classes
- Targets can be:
 - arbitrary integers / names
 - one-hot encoding
 - N vectors ; each with one element = 1, others = 0
- Can use binary classifiers for multiclass problems
 - One-vs-Rest strategy (default)
 - One-vs-One strategy
 - Scikit-learn classifiers do this automatically



Confusion Matrix

- Displays predicted vs actual class
 - run on the test set
- Diagonal elements are good
- Off-diagonals show errors
- Can have any number of classes
- Shows where most common errors occur

Predicted: Actual:	Class 1 (Negative)	Class 2 (Positive)
Class 1	True Negative count	False Positive count
Class 2	False Negative count	True Positive count

```
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import confusion_matrix
y_train_pred = cross_val_predict(clf, X_train, y_train)
confusion_matrix(y_train, y_train_pred)
```

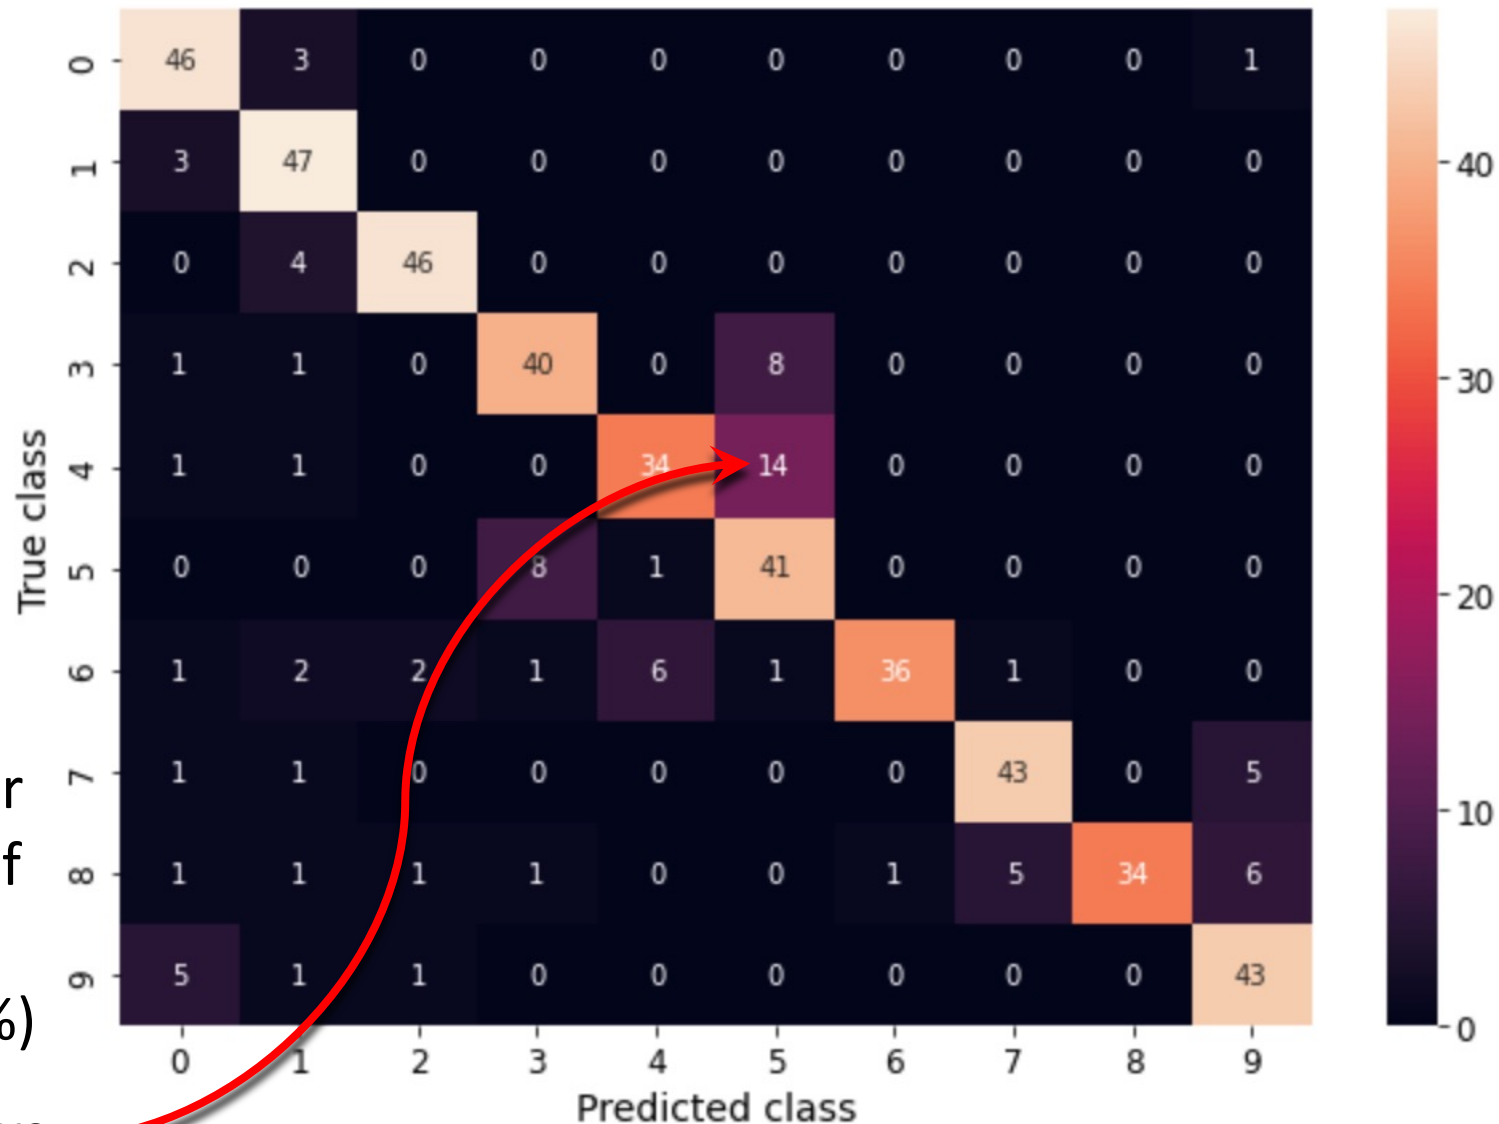
Confusion Matrix

Sum within a row =
number of true
samples of that class

Sum within col =
number of predictions
of that class

Can show raw counts or
normalised (e.g. as % of
true class samples so
each row sums to 100%)

Large off-diagonal shows
common error: here the true
class 4 is predicted as class 5



```
import seaborn as sn
sn.heatmap(array, annot=True)
```

Summary

- Classification vs. regression
- Range of classifier approaches
 - SGD classifier
 - Decision Tree
- White box vs. Black box
 - Decision boundaries
- Several performance metrics
 - Precision, recall
 - ROC, AUC
 - Confusion matrix
- Multi-class classification
 - Different algorithms for combining 2-class classification
 - Different options for representing labels (esp. one-hot encoding)